

Encoding Separation Logic in SMT-LIB v2.5

Radu Iosif², Cristina Serban², Andrew Reynolds¹, and Mihaela Sighireanu³

¹ The University of Iowa

² Verimag/CNRS/Université de Grenoble Alpes

³ IRIF/Université Paris Diderot

Abstract. We propose an encoding of Separation Logic using SMT-LIB v2.5. This format is currently supported by SMT solvers (CVC4) and inductive proof-theoretic solvers (SLIDE and SPEN). Moreover, we provide a library of benchmarks written using this format, which stems from the set of benchmarks used in SL-COMP'14 [?].

1 Preliminaries

We consider formulae in multi-sorted first-order logic. A *signature* Σ consists of a set Σ^s of sort symbols and a set Σ^f of *function symbols* $f^{\sigma_1 \dots \sigma_n \sigma}$, where $n \geq 0$ and $\sigma_1, \dots, \sigma_n, \sigma \in \Sigma^s$. If $n = 0$, we call f^σ a *constant symbol*. We make the following assumptions:

1. all signatures Σ contain the Boolean sort \mathbf{B} , where \top and \perp denote the Boolean constants *true* and *false*.
2. Σ^f contains a boolean equality function $\approx^{\sigma \mathbf{B}}$ for each sort symbol $\sigma \in \Sigma^s$.

Let \mathbf{Vars} be a countable set of first-order variables, each $x^\sigma \in \mathbf{Vars}$ having an associated sort σ . First-order terms and formulae over the signature Σ (called Σ -terms and Σ -formulae) are defined as usual. A Σ -*interpretation* \mathcal{I} maps:

- each sort symbol $\sigma \in \Sigma$ to a non-empty set $\sigma^\mathcal{I}$,
- each function symbol $f^{\sigma_1 \dots \sigma_n \sigma} \in \Sigma$ to a total function $f^\mathcal{I} : \sigma_1^\mathcal{I} \times \dots \times \sigma_n^\mathcal{I} \rightarrow \sigma^\mathcal{I}$ where $n > 0$, and to an element of $\sigma^\mathcal{I}$ when $n = 0$, and
- each variable $x^\sigma \in \mathbf{Vars}$ to an element of $\sigma^\mathcal{I}$.

For an interpretation \mathcal{I} a sort symbol σ and a variable x , we denote by $\mathcal{I}[\sigma \leftarrow S]$ and, respectively $\mathcal{I}[x \leftarrow v]$, the interpretation associating the set S to σ , respectively the value v to x , and which behaves like \mathcal{I} in all other cases. By writing $\mathcal{I}[\sigma \leftarrow S]$ we ensure that all variables of sort σ are mapped by \mathcal{I} to elements of S . For a Σ -term t , we write $t^\mathcal{I}$ to denote the interpretation of t in \mathcal{I} , defined inductively, as usual. A satisfiability relation between Σ -interpretations and Σ -formulas, written $\mathcal{I} \models \varphi$, is also defined inductively, as usual. In this case, we say that \mathcal{I} is a *model* of φ .

A (multi-sorted first-order) *theory* is a pair $T = (\Sigma, \mathbf{I})$ where Σ is a signature and \mathbf{I} is a non-empty set of Σ -interpretations, the *models* of T . A Σ -formula φ is *T-satisfiable* if it is satisfied by some interpretation in \mathbf{I} .

2 Ground Separation Logic

Let $T = (\Sigma, \mathbf{I})$ be a theory and let \mathbf{Loc} and \mathbf{Data} be two sorts from Σ , with no restriction other than the fact that \mathbf{Loc} is always interpreted as a countable set. Also, we consider

that Σ has a designated constant symbol nil^{Loc} . We define the *Ground Separation Logic* $\text{SL}(T)_{\text{Loc,Data}}^g$ to be the set of formulae generated by the following syntax:

$$\varphi := \phi \mid \text{emp} \mid t \mapsto u \mid \varphi_1 * \varphi_2 \mid \varphi_1 \multimap \varphi_2 \mid \neg \varphi_1 \mid \varphi_1 \wedge \varphi_2 \mid \exists x^\sigma . \varphi_1(x)$$

where ϕ is a Σ -formula, and t, u are Σ -terms of sorts Loc and Data , respectively. As usual, we write $\forall x^\sigma . \varphi(x)$ for $\neg \exists x^\sigma . \neg \varphi(x)$. We omit specifying the sorts of variables and functions when they are clear from the context.

Given an interpretation \mathcal{I} , a *heap* is a finite partial mapping $h : \text{Loc}^{\mathcal{I}} \rightarrow_{\text{fin}} \text{Data}^{\mathcal{I}}$. For a heap h , we denote by $\text{dom}(h)$ its domain. For two heaps h_1 and h_2 , we write $h_1 \# h_2$ for $\text{dom}(h_1) \cap \text{dom}(h_2) = \emptyset$ and $h = h_1 \uplus h_2$ for $h_1 \# h_2$ and $h = h_1 \cup h_2$. We define the *satisfaction relation* $\mathcal{I}, h \models_{\text{SL}} \phi$ inductively, as follows:

$$\begin{aligned} \mathcal{I}, h \models_{\text{SL}} \phi &\iff \mathcal{I} \models \phi \text{ if } \phi \text{ is a } \Sigma\text{-formula} \\ \mathcal{I}, h \models_{\text{SL}} \text{emp} &\iff h = \emptyset \\ \mathcal{I}, h \models_{\text{SL}} t \mapsto u &\iff h = \{(t^{\mathcal{I}}, u^{\mathcal{I}})\} \text{ and } t^{\mathcal{I}} \neq \text{nil}^{\mathcal{I}} \\ \mathcal{I}, h \models_{\text{SL}} \phi_1 * \phi_2 &\iff \text{there exist heaps } h_1, h_2 \text{ s.t. } h = h_1 \uplus h_2 \text{ and } \mathcal{I}, h_i \models_{\text{SL}} \phi_i, i = 1, 2 \\ \mathcal{I}, h \models_{\text{SL}} \phi_1 \multimap \phi_2 &\iff \text{for all heaps } h' \text{ if } h' \# h \text{ and } \mathcal{I}, h' \models_{\text{SL}} \phi_1 \text{ then } \mathcal{I}, h' \uplus h \models_{\text{SL}} \phi_2 \\ \mathcal{I}, h \models_{\text{SL}} \exists x^S . \varphi(x) &\iff \mathcal{I}[x \leftarrow s], h \models_{\text{SL}} \varphi(x), \text{ for some } s \in S^{\mathcal{I}} \end{aligned}$$

The satisfaction relation for Σ -formulae, Boolean connectives \wedge, \neg , and linear arithmetic atoms, are the classical ones from first-order logic. Notice that the range of a quantified variable x^S is the interpretation of its associated sort $S^{\mathcal{I}}$. A formula φ is said to be *satisfiable* if there exists an interpretation \mathcal{I} and a heap h such that $\mathcal{I}, h \models_{\text{SL}} \varphi$. We say that φ *entails* ψ , written $\varphi \models_{\text{SL}} \psi$, when every pair (\mathcal{I}, h) which satisfies φ , also satisfies ψ .

2.1 SMT-LIB Encoding

We write ground SL formulae in SMT-LIB using the following functions:

```
(emp Bool)
(sep Bool Bool Bool :left-assoc)
(wand Bool Bool Bool :right-assoc)
(par (Loc Data) (pto Loc Data Bool))
(par (Loc) (nil Loc))
```

Observe that `pto` and `nil` are polymorphic functions, with sort parameters `Loc` and `Data`. There is no restriction on the choice of `Loc` and `Data`, as shown below.



Is `nil` always of sort `Loc` ? Here `Loc` is just a parameter with no special meaning.

Assume that `Loc` is an uninterpreted sort `U`, `Data` is the integer sort `Int`, and that $x^{\text{U}}, y^{\text{U}}, a^{\text{Int}}$ and b^{Int} are constants:

```
(declare-sort U 0)
```

```
(declare-const x U)
(declare-const y U)
(declare-const a Int)
(declare-const b Int)
```

We write the SL formula $\text{emp} \wedge ((x \mapsto a * y \mapsto b) \multimap (x \mapsto \text{nil} * \top))$ in SMT-LIB as follows:

```
(and emp (wand (sep (pto x a) (pto y b)) (sep (pto x (as nil Int)) true))))
```

In addition to the classical SMT-LIB typing constraints, the SL theories require that the heap models are well-typed. For instance, a separation constraint of the form:

```
(sep (pto x y) (pto a b))
```

with the above constant declarations results in a typing error, because $(\text{pto } x \ y)$ requires the heap to be of type $U \rightarrow U$, whereas $(\text{pto } a \ b)$ requires the heap to be of type $\text{Int} \rightarrow \text{Int}$, and combining heaps of different types is not allowed.

1. It is currently unclear what the type of the heap introduced by `emp` should be. The solution currently adopted in CVC4 is to give `emp` parameters used only to infer the type, as in `(emp x a)`. Another solution, used in SLIDE, is to infer the type of `emp` based on the context. For instance the type of the `emp` heap in `(sep (pto x a) emp)` is $U \rightarrow \text{Int}$. We leave this point for discussion.
2. It is usually expected that `nil` be of type `Loc`, however currently one can force other type than `Loc`, like in the previous example. Shall we impose a stricter type checking on `nil` ?

This heap typing restriction is not a limitation of the expressive power of the SMT-LIB encoding and can be easily overcome by using datatypes (available in SMT-LIB v2.5). Suppose, for instance that we would like to specify a heap consisting of cells containing both integer and boolean data. The idea is to declare a union type:

```
(declare-datatype BoolInt ((cons_bool (d Bool)) (cons_int (d Int))))
```

and use it to describe a mixed data heap, as in:

```
(sep (pto x (cons_bool false)) (pto y (cons_int 0)))
```

The same workaround can be used to specify heaps with mixed addresses, although this is a much less common situation in practice.

2.2 Separation Logic with Inductive Definitions

Let Pred be a set of second-order variables, each $P^{\sigma_1 \dots \sigma_n} \in \text{Pred}$ having an associated tuple of parameter sorts $\sigma_1, \dots, \sigma_n \in \Sigma^s$. In addition to the first-order terms built using variables from Vars and function symbols from Σ^f , we enrich the language of SL with the boolean terms $P^{\sigma_1 \dots \sigma_n}(t_1, \dots, t_n)$, where each t_i is a first-order term of sort σ_i , for $i = 1, \dots, n$. Each second-order variable $P^{\sigma_1 \dots \sigma_n} \in \text{Pred}$ is provided with an inductive definition $P(x_1, \dots, x_n) \leftarrow \phi_P(x_1, \dots, x_n)$, where ϕ_P is a formula in the extended language, possibly containing occurrences of P . The satisfaction relation is then extended as follows:

$$\mathcal{I}, h \models_{\text{SL}} P^{\sigma_1 \dots \sigma_n}(t_1, \dots, t_n) \iff \mathcal{I}, h \models_{\text{SL}} \phi_P(t_1^I, \dots, t_n^I)$$

where ϕ_P is the inductive definition of $P^{\sigma_1 \dots \sigma_n}$. Observe that, given a set of inductive definitions, the set of possible models for each second-order variable is the least fixed point of a monotonic and continuous function mapping tuples of sets of models to a set of models.

2.3 SMT-LIB Encoding

An inductive definition $P(x_1, \dots, x_n) \leftarrow \phi_P(x_1, \dots, x_n)$ is written in SMT-LIB using a recursive function definition. For instance, the inductive definition of a doubly-linked list segment:

$$\text{dllseg}(h, p, t, n) \leftarrow (\text{emp} \wedge h \approx n \wedge p \approx t) \vee (\exists x^{\text{Loc}}. h \mapsto (x, p) * \text{dllseg}(x, h, t, n))$$

is written into SMT-LIB as follows:

```
(declare-datatype Node ((node (next Loc) (prev Loc))))

(define-fun-rec dllseg ((h Loc) (p Loc) (t Loc) (n Loc)) Bool
  (or (and emp (= h n) (= p t))
      (exists ((x Loc)) (sep (pto h (node x p)) (dllseg x h t n)))
  )
)
```

2.4 A Detailed Example

Let us go through an example step by step. First of all, the preamble of and SMT-LIB file describing a SL satisfiability query must contain (at least):

```
(set-logic SEPLUG)
```

Since SL is used in combination with other theories, it is customary to start with:

```
(set-logic ALL_SUPPORTED)
```

We consider the slightly modified version of the dllseg definition above, which describes a doubly-linked list segment with ordered integer data:

$$\text{dllseg}_{\text{ord}}(h, p, t, n, \text{min}) \leftarrow (\text{emp} \wedge h \approx n \wedge p \approx t) \vee (\exists x^{\text{Loc}} \exists d^{\text{Int}} . h \mapsto (d, x, \text{min}) * \text{dllseg}_{\text{ord}}(x, h, t, n, d)) \wedge \text{min} \leq d$$

Since we do not perform any pointer arithmetic reasoning, we can declare Loc to be an uninterpreted sort:

```
(declare-sort U 0)
```

We encode the definition of $\text{dllseg}_{\text{ord}}$ as:

```
(declare-datatype Node ((node (data Int) (next Loc) (prev Loc))))

(define-fun-rec dllseg_ord ((h Loc) (p Loc) (t Loc) (n Loc) (min Int)) Bool
  (or (and emp (= h n) (= p t))
      (exists ((x Loc) (d Int)) (and
        (sep (pto h (node x p)) (dllseg_ord x h t n))
        (<= min d)
      )
    )
  )
)
```

Let us consider the problem of proving that a $\text{dllseg}_{\text{ord}}$ to which a node is appended is again a $\text{dllseg}_{\text{ord}}$, provided that the data of the new node is smaller than the minimal element of the first $\text{dllseg}_{\text{ord}}$:

$$x \mapsto (m, u, v) * \text{dllseg}_{\text{ord}}(u, x, z, t, n) \wedge m \leq n \models_{\text{SL}} \text{dllseg}_{\text{ord}}(x, y, z, t, m)$$

We encode this entailment problem as an assertion asking whether the negated problem is satisfiable:

```
(declare-const x U)
(declare-const y U)
(declare-const z U)
(declare-const u U)
(declare-const v U)
(declare-const t U)
(declare-const m Int)
(declare-const n Int)

(assert (not (implies
  (and (sep (pto x (node m u v)) (dllseg_ord u x z t n)) (<= m n))
  (dllseg_ord x y z t m)
)
))
)
```

The entailment holds when the assertion is unsatisfiable, which can be checked in the standard way, using `(check-sat)`. However, the dual problem:

```
(assert (not (implies
  (dllseg_ord x y z t m)
  (and (sep (pto x (node m u v)) (dllseg_ord u x z t n)) (<= m n))
)))
)
```

is satisfiable, and the counter-model can be obtained in the standard way, using `(get-model)`. Observe that the model of a satisfiable SL query consists of an interpretation of the constants and a specification of the heap.

3 Abduction and Frame Inference

Abduction and frame inference (or bi-abduction for both) are problems that occur in the context of program verification. In this case, the solver is not only required to give a yes/no answer to a satisfiability query, but to infer SL formulae that ensure the validity of a given entailment. Given SL formulae φ and ψ , and second-order variables X and Y , we consider the following synthesis problems:



TO BE REFINED: what are the first-order parameters of X and Y ? It seems that X is allowed to all free variables of φ , but how about Y ? Is it allowed to all free variables of φ and ψ , or just the ones of ψ ?

1. The *abduction problem* asks for a satisfiable definition of a X such that $\varphi * X \models_{\text{SL}} \psi$. Sometimes X is called an *anti-frame*. Observe that $X \leftarrow \perp$ is always a solution, but not a very interesting one.
2. The *frame inference problem* asks for a definition of Y such that $\varphi \models_{\text{SL}} \psi * Y$.
3. The *bi-abduction problem* asks for both a satisfiable definition of X and a definition of Y such that $\varphi * X \models_{\text{SL}} \psi * Y$.

The capability of solving the above problems is key to using a given SL solver for practical program verification purposes. For this reason, we aim at finding a standard way of specifying these problems in SMT-LIB.

4 Additional Resources

The quest for a suitable format for SL solvers started with SL-COMP'14 [?], which adopted the QF_S format, described in [?]. The current proposal is inspired by QF_S, and relies on the datatypes introduced SMT-LIB v2.5 for an elegant encoding of union and record types. The tools supporting SMT-LIB as a native language are:

- CVC4 [?] – a description of the SL format of CVC4 is provided in [?] (a slightly modified version of the current proposal)

- SLIDE (under construction) – uses the encoding from the current proposal.
- SPEN [?] – a description of the SL format of SPEN (QF_S) is available in [?].

Other tools that participated to SL-COMP'14 have been adapted to QF_S by means of a specialized front-end [?]. It is our goal to convince the developpers of SL solvers to adopt SMT-LIB as the native input language of their tools. For this purpose, we provide a C++ front-end [?] that can be used to parse and type check SL inputs encoded in SMT-LIB using the current specification.